

## **Stock Market Metrics Computation Architecture – Complete AWS**

This architecture provides a robust, scalable solution for real-time and historical stock market data processing with strategy evaluation across 500+ trading strategies using 1-minute candle data. The system ingests market data, processes it through computational phases, evaluates strategies, and delivers alerts/orders through a notification service network.

It is a layered architecture and each layer can be treated as a subnet within custom vpc.

Key Components:

### **1. Data Ingestion Layer:**

- WebSocket connections for real-time market data streaming into Kinesis Data Streams
- Amazon S3 for historical data storage (CSV, PDFs)
- AWS Certificate Manager for SSL/TLS security

### **2. Computation Layer:**

- AWS Fargate for containerized strategy evaluation (scalable and serverless)
- AWS Batch for historical data processing
- Aurora PostgreSQL (Serverless) as sync storage for strategy state

### **3. Notification & Order Processing:**

- SNS for alert distribution with topics
- SQS queues (Order Alerts and Webhook Alerts) for reliable message delivery
- Lambda functions for broker API integration

#### **4. Operational Excellence:**

- CloudWatch for monitoring and logging
- X-Ray for distributed tracing and debugging
- VPC for network isolation and security

#### **5. Infrastructure as Code:**

- Terraform for provisioning and lifecycle management
- IGW for controlled internet access

### **IaC Terraform Implementation Strategy :**

#### **Modules:**

##### **1) Networking**

- VPC  
Security Groups
- Load Balancers

##### **2) Data Ingestion**

- Kinesis
- S3

##### **3) Compute**

- Fargate
- Lambda
- Batch

##### **4) Database**

- Aurora with Postgres

## **5)Notification**

- SNS
- SQS

## **6)Monitoring**

- Cloudwatch
- Xray

## **7)Security**

- IAM
- Cert-manager

## **Environment Strategy**

### **1. Workspace Approach:**

- dev, staging, prod workspaces for environment isolation
- Shared modules with environment-specific variables

### **2. State Management:**

- Remote state in S3 with DynamoDB locking
- Separate state files per environment
- State isolation between modules where appropriate

### **3. Variable Strategy:**

- terraform.tfvars for common variables
- dev.tfvars, staging.tfvars, prod.tfvars for environment specifics

- **Sensitive variables in AWS Secrets Manager with Terraform data sources**

## **Technology Explained:**

### **1. AWS Fargate over ECS/EKS:**

- Serverless operation reduces management overhead
- Automatic scaling matches strategy evaluation workload
- Cost-effective for bursty computation patterns

### **2. Aurora PostgreSQL Serverless:**

- Auto-scaling matches trading session patterns
- Cost optimization during off-market hours
- PostgreSQL's advanced analytical capabilities

### **3. Kinesis over Kafka:**

- Fully managed service reduces operational burden
- Seamless integration with AWS analytics services
- Built-in scaling for variable market data volumes

### **4. SNS/SQS over direct API calls:**

- Decouples strategy evaluation from broker integration
- Provides retry mechanisms and dead-letter handling
- Enables multiple alert consumers without modification

### **5. Terraform over CloudFormation/CDK:**

- Multi-cloud potential for future expansion

- Rich module ecosystem for financial services patterns
- Better state management for complex infrastructures

## Security, Scalability, and Availability

### Security Implementation

#### 1. Data Protection:

- TLS everywhere (Certificate Manager)
- VPC isolation with security groups limiting east-west traffic
- IAM roles with least privilege (per-service roles)
- Secrets management via AWS Secrets Manager

#### 2. Network Security:

- VPC endpoints for AWS services to avoid public internet
- Web Application Firewall (WAF) on load balancers
- Security groups with minimum necessary ports

#### 3. Operational Security:

- CloudTrail enabled for all API calls
- Config Rules for compliance monitoring
- Regular security scans of container images

### Scalability Patterns

#### 1. Horizontal Scaling:

- Fargate services auto-scale based on SQS queue depth
- Kinesis shards adjust based on incoming data rate

- Aurora Serverless scales compute based on demand

## **2. Decoupled Architecture:**

- SQS buffers between computation and notification
- Lambda concurrency limits prevent broker API overload
- Batch job arrays for parallel historical processing

## **3. Performance Optimization:**

- Read replicas for Aurora during heavy analysis
- Data partitioning in S3 by date/symbol
- Caching layer potential (ElastiCache) for frequent indicators

## Availability Design

### **1. Multi-AZ Deployment:**

- Aurora across 3 AZs
- Fargate tasks distributed across AZs
- S3 with cross-region replication for critical data

### **2. Fault Tolerance:**

- Dead-letter queues for failed messages
- Circuit breakers in Lambda functions
- Retry logic with exponential backoff

### **3. Disaster Recovery:**

- Terraform modules support multi-region deployment
- Regular S3 backups with versioning

- Automated failover testing procedures

## **Modules:**

### **Networking:**

```
module "vpc" {  
    source = "terraform-aws-modules/vpc/aws"  
    version = ""  
  
    name = "trading-vpc"  
    cidr = var.vpc_cidr  
    azs      = var.availability_zones  
    private_subnets = var.private_subnets  
    public_subnets  = var.public_subnets  
  
    enable_nat_gateway = true  
    single_nat_gateway = false  
}
```

### **Fargate Service:**

```
module "strategy_evaluator" {  
    source = "../modules/compute/fargate"  
  
    cluster_name    = "strategy-evaluation"  
    service_name    = "evaluator"  
    task_cpu        = 4096  
    task_memory     = 8192  
    container_image = var.evaluator_image  
    desired_count   = var.environment == "prod" ? 3 : 1  
  
    vpc_id          = module.vpc.vpc_id  
    subnets        = module.vpc.private_subnets  
    security_group_ids =  
[module.sg_strategy_evaluator.security_group_id]  
  
    sqs_queue_arns    = [module.order_alerts.queue_arn]  
    db_secret_arn      =  
aws_secretsmanager_secret_version.db_credentials.arn  
}
```

### **Notification Pipeline:**

```
module "trading_alerts" {
```



```
source = "./modules/notifications"
```

```
alert_topics = {  
    "high-priority" = { protocol = "sqs" },  
    "medium-priority" = { protocol = "lambda" },  
    "low-priority" = { protocol = "email" }  
}
```

```
lambda_config = {  
    runtime      = "python3.12"  
    handler      = "alert_processor.lambda_handler"  
    source_dir   = "../src/lambda/alert_processor"  
    memory_size  = 512  
    timeout      = 30  
}  
}
```